

236606 - Deep Learning - Homework 3

Eyal Ben David - 305057416

June 5, 2018

Convolutional Neural Network


During our assignment we were asked to build a cifar10 classifier with a maximum of 50K parameters, and a min accuracy of - 85%. We tried multiple types of architecture's in order to meet with these requirements including; very deep (9-11 layers) convolutional networks with shallow channels (8-16), shallow convolutional layers (3-4 layers) with wide channels(48-64) proceeded by wide dense layers and . We also explored different optimizers such as adam, rmsprop and sgd and different regularization technique like L1, L2, dropout and data augmentation. Our winning architecture is a fully convolutional architecture with six convolutional layers, with 'ReLU' activation and increasing channel width succeeded by a softmax layer. Here is our model summary:

Figure 1: Model summary of our chosen model

Layer (type)	Output Shape	Param #
conv2d_1 (Conv2D)	(None, 32, 32, 15)	420
batch_normalization_1 (Batch Normalization)	(None, 32, 32, 15)	60
conv2d_2 (Conv2D)	(None, 32, 32, 15)	2040
batch_normalization_2 (Batch Normalization)	(None, 32, 32, 15)	60
max_pooling2d_1 (MaxPooling2D)	(None, 16, 16, 15)	0
conv2d_3 (Conv2D)	(None, 16, 16, 28)	3808
batch_normalization_3 (Batch Normalization)	(None, 16, 16, 28)	112
conv2d_4 (Conv2D)	(None, 16, 16, 28)	7084

batch_normalization_4	(Batch (None, 16, 16, 28))	112
max_pooling2d_2	(MaxPooling2 (None, 8, 8, 28))	0
conv2d_5	(Conv2D) (None, 8, 8, 42)	10626
batch_normalization_5	(Batch (None, 8, 8, 42))	168
conv2d_6	(Conv2D) (None, 8, 8, 42)	15918
batch_normalization_6	(Batch (None, 8, 8, 42))	168
max_pooling2d_3	(MaxPooling2 (None, 4, 4, 42))	0
flatten_1	(Flatten) (None, 672)	0
dropout_1	(Dropout) (None, 672)	0
batch_normalization_7	(Batch (None, 672))	2688
dense_1	(Dense) (None, 10)	6730
=====		
Total params: 49,994		
Trainable params: 48,310		
Non-trainable params: 1,684		

Figure 2: Model visualization of our chosen model




To train this model we used the 'adam' optimizer with Keras default parameters

for 400 epochs. Another technique from keras which gave us a good boost for the accuracy was the 'Data Generator', which we used in order to augment the data and decrease overfitting of our CNN. The final results of this model are presented in 'table 1' and the training graph is plotted in 'figure 2'. The trained model is submitted together with the code in the .zip file.

Table 1: Chosen model results for 'Cifar10' dataset

	Loss	Accuracy
Train	0.30948	0.8922
Test	0.44967	0.8536

Figure 2: Training accuracy graph of the model in each epoch



Transfer Learning

In this question we have investigated some transfer learning approaches that include: 'Fine Tuning', 'KNN', 'SVM' (our flavor). The original model is a VGG16 model trained on 'Cifar100' dataset. For adapting to the 'Cifar10' dataset we popped the last layer of the original model, and for each section we added a new layer. The number of examples was hardcoded inside the file. For each approach we have trained the model with 100, 1000 and 10000 examples.

Transfer learning by Fine Tuning

We took a CIFAR-100 trained model and replaced the last fully-connected layer with a new fully-connected layer that is adapted to ten classes output vector. We then trained only the last layer using 100, 1000, and 10000 examples on CIFAR-10 dataset, while freezing other layers of the model.

Transfer learning with KNN

In this approach we map the new CIFAR-10 dataset to the domain that was trained by the CIFAR-100 dataset. After passing the data in the pre-trained model we get vectors in dimension of the last output layer - 512, then we apply KNN to the 5 closest neighbours. This can work because the classes of CIFAR-10 and CIFAR-100 have a strong correlation between them. We used 5 closest neighbours because it gave the best results of all other 'K' options we have tried.

Transfer learning with SVM

In this section we needed to choose our own solution, we chose to use SVM. Similarly to the KNN method, we use the pretrained model of CIFAR-100 dataset, without its last layer. We pass the new CIFAR-10 data through the model and get a new representation for each example. Then we train these new vectors with a non-linear classifier (we are using gaussian kernel) to classify 10 different classes. In each stage we trained the non-linear classifier with different amount of examples. The motivation for this approach is that if the previous layers did a good job embedding the data to a new dimension which the data is separated in this dimension, then it might be classified easily by the non-linear classifier. SVM can also use other kernels in order to use a non-linear classifier for the data after it is projected on the new dimension.

Results - Transfer learning

We summarize the results of this question in 'table 2'.

Table 2: Accuracy for each Transfer Learning model trained on different number of examples

	Fine Tune Train	Fine Tune Test	KNN (neighbours = 5)	SVM
100 examples	70.3%	49.25%	43.15%	48.35%
1,000 examples	56.5%	56.2%	50.88%	57.01%
10,000 examples	61.2%	59.1%	57.51%	62.41%

Conclusions

From these results we can learn a couple of things: The more examples we use to train, the accuracy is higher and overfitting decreases. Also, the 'SVM' and 'Fine Tuning' methods work better for transfer learning with this model than 'KNN' method. For each method we trained only the last layer of the model, while we froze other layers in the model that were trained to classify the CIFAR-100 dataset. Hence it is clear that the more correlation we get between the two datasets, their types of classes, the classes distribution and the similarity of the data itself, the better the results can get. We see that the results are pretty good

in the case of using a pretrained CIFAR-100 model on a CIFAR-10 dataset, and that clarifies that the two dataset are correlated.